

# Class 7: Machine Learning 1

Eric Wang (PID: A17678188)

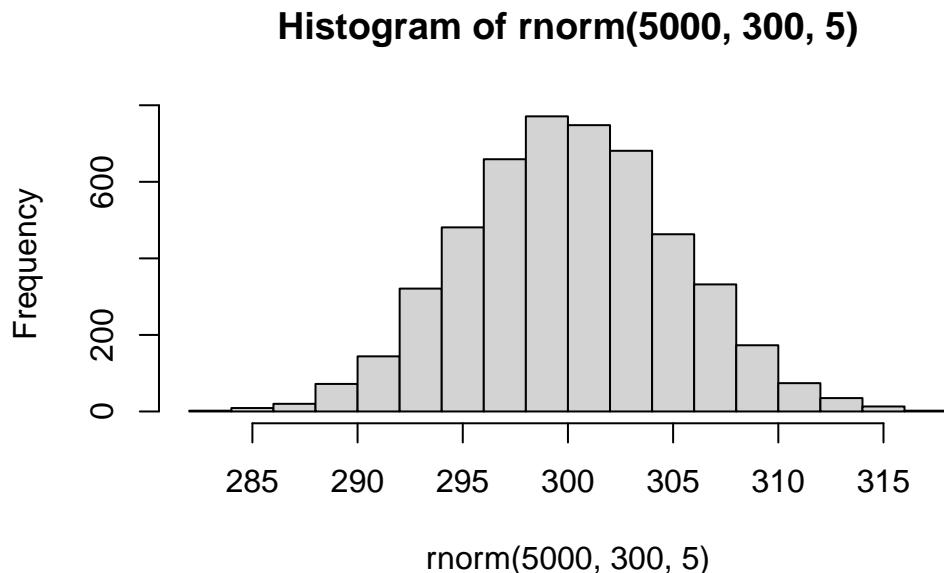
## Background

Today we will begin our exploration of some important machine learning methods, namely **clustering** and **dimensionality reduction**.

Let's make up some input data for clustering where we know what the natural "clusters" are.

The function `rnorm()` can be useful here.

```
hist(rnorm(5000, 300, 5))
```



Q. Generate 30 random numbers centered at + 3 and -3

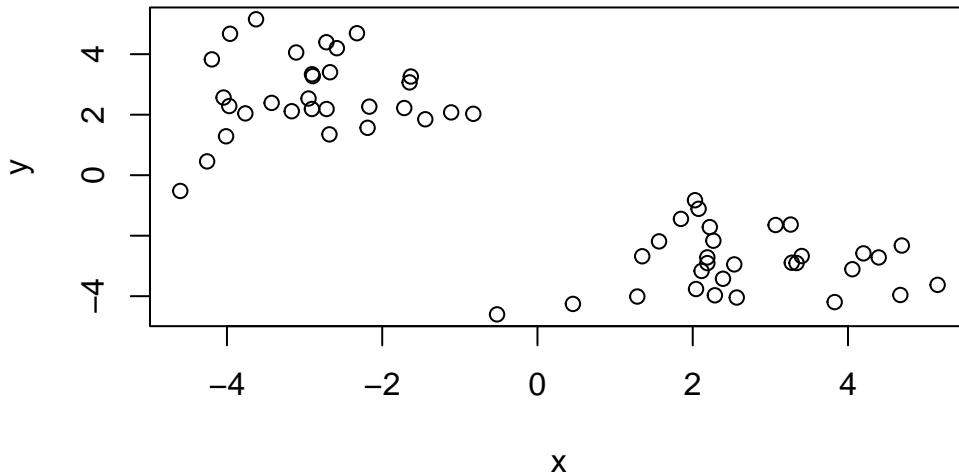
```

tmp <- c(rnorm(30, 3), rnorm(30,-3))

x <- cbind(x= tmp, y=rev(tmp))

plot(x)

```



## K-means clustering

The main function in “base R” for K-means clustering is called `kmeans()` :

```

km <- kmeans(x, 2)
km

```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.873557	2.674594
2	2.674594	-2.873557

Clustering vector:

```
Within cluster sum of squares by cluster:  
[1] 76.60792 76.60792  
(between_SS / total_SS =  85.8 %)
```

### Available components:

```
[1] "cluster"        "centers"        "totss"          "withinss"        "tot.withinss"  
[6] "betweenss"     "size"           "iter"           "ifault"
```

Q. What component of the results object details the cluster sizes?

km\$size

[1] 30 30

Q. What component of the results object details the cluster centers?

km\$centers

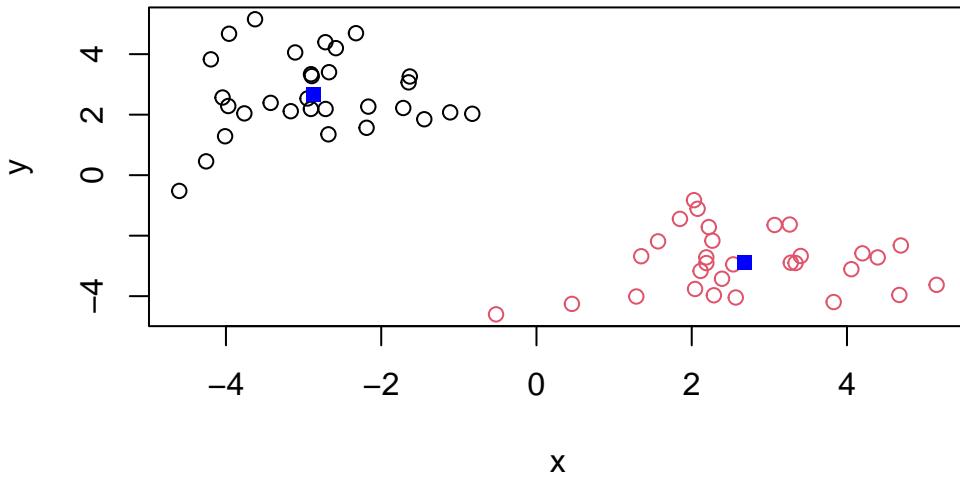
	x	y
1	-2.873557	2.674594
2	2.674594	-2.873557

Q. What component of the results object details the cluster membership vector (i.e. our main result of which points lie in which cluster)

km\$cluster

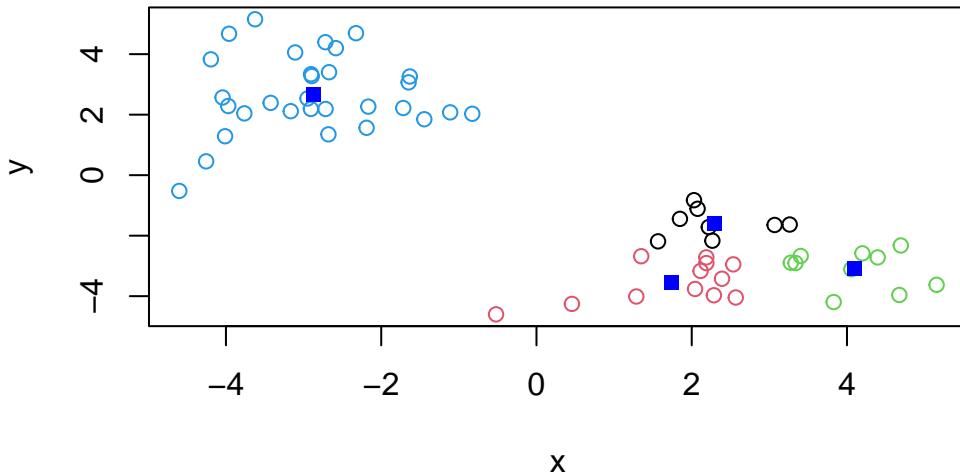
Q. Plot our clustering results with points colored by cluster and also add the cluster centers as new points colored blue?

```
plot(x,col=km$cluster)
points(km$centers, col="blue",pch=15)
```



Q. Run `kmeans()` again and this time produce 4 clusters and call your result object `k4`

```
k4 <- kmeans(x, 4)
plot(x, col= k4$cluster)
points(k4$centers, col="blue", pch=15)
```



The metric

```
k4$tot.withinss
```

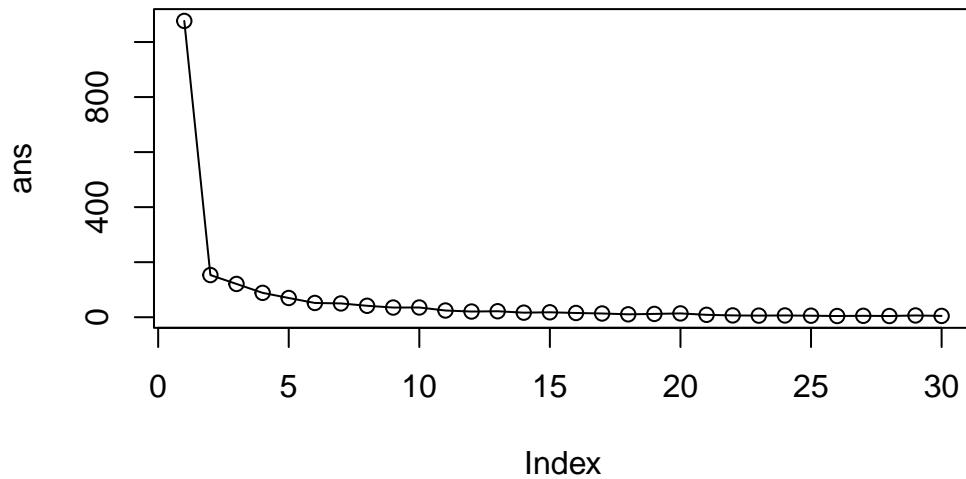
```
[1] 102.1992
```

Q. Let's try different number of K(centers) from 1 to 30 and see what the best result is?

```
ans <- NULL
for (i in 1:30){
  ans <- c(ans, kmeans(x,centers = i)$tot.withinss)
}
ans
```

```
[1] 1076.675070 153.215848 120.890409 88.564969 69.873721 51.790720
[7] 50.056354 41.547306 35.097239 35.334373 24.402346 20.612429
[13] 21.670084 16.574259 17.906721 15.364529 13.516267 10.500297
[19] 11.941704 13.868758 8.708025 6.595800 5.822908 6.368789
[25] 5.502479 4.368836 5.140008 4.295194 6.359370 4.731376
```

```
plot(ans, typ= "o")
```



`tot.withnss` shows how tight the cluster it is. The lower the value the tighter the clusters group.

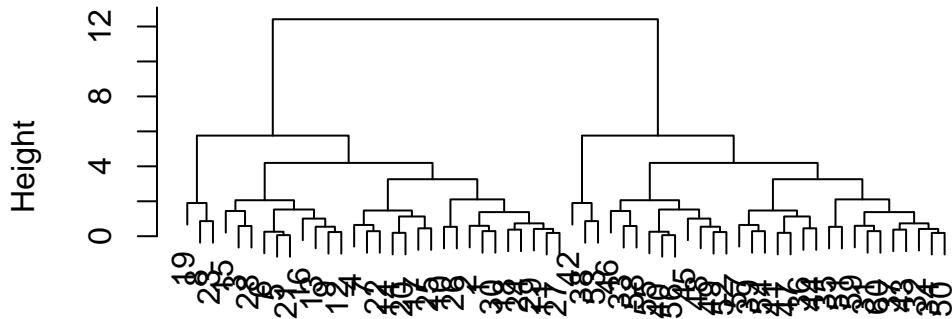
**Key-Point:** K-means will impose a clustering structure on your data even if it is not there - it will always give you the answer you asked for even if that answer is silly!

## Hierarchical Clustering

The main function for Hierarchical Clustering is called `hclust()`. Unlike `kmeans()` (which does all the work for you) you can't just pass `hclust()` our raw input data. It needs a "distance matrix" like the one returned from the `dist()` function.

```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```

## Cluster Dendrogram



```
d  
hclust (*, "complete")
```

To extract our cluster membership vector from a `hclust()` result object we have to “cut” our tree at a given height to yield separate “groups”/“branches”.

To do this we use the `cutree()` function on our `hclust()` objection:

```
grps <- cutree(hc, h=8)  
grps
```

```
table(grps, km$cluster)
```

grps	1	2
1	0	30
2	30	0

## PCA of UK food data

Import the data set of food consumption in the UK

```

url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
x

```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139
7	Fresh_potatoes		720	874	566	1033
8	Fresh_Veg		253	265	171	143
9	Other_Veg		488	570	418	355
10	Processed_potatoes		198	203	220	187
11	Processed_Veg		360	365	337	334
12	Fresh_fruit		1102	1137	957	674
13	Cereals		1472	1582	1462	1494
14	Beverages		57	73	53	47
15	Soft_drinks		1374	1256	1572	1506
16	Alcoholic_drinks		375	475	458	135
17	Confectionery		54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

One solution to set the row names is by hand...

```

#rownames(x)
rownames(x) <- x[,1]

```

To remove the first column I can use the minus index trick

```

x <- x[,-1]
x

```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

A better way to do this is to set the row names to the first column with `read.csv()`

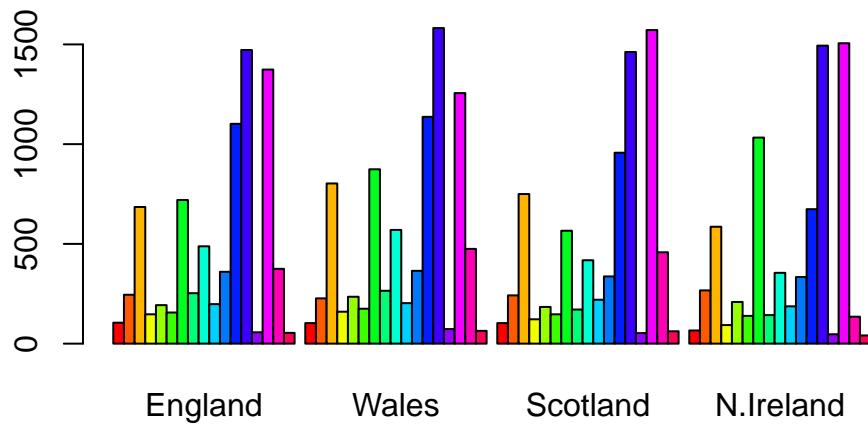
```
x <- read.csv(url, row.names = 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

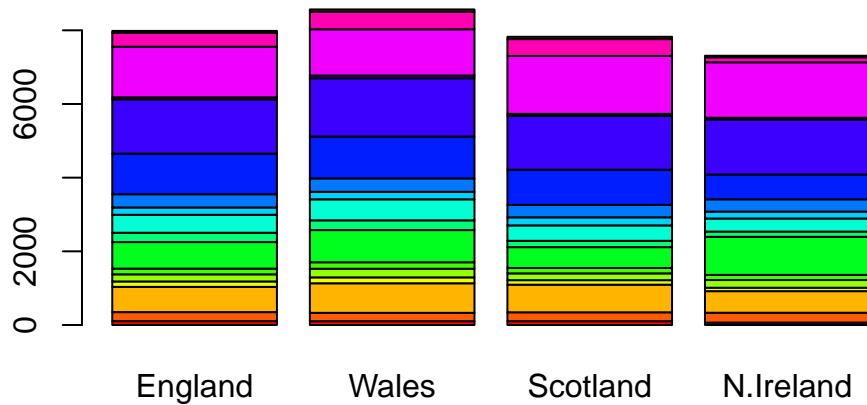
Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

### Spotting major differences and trends

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

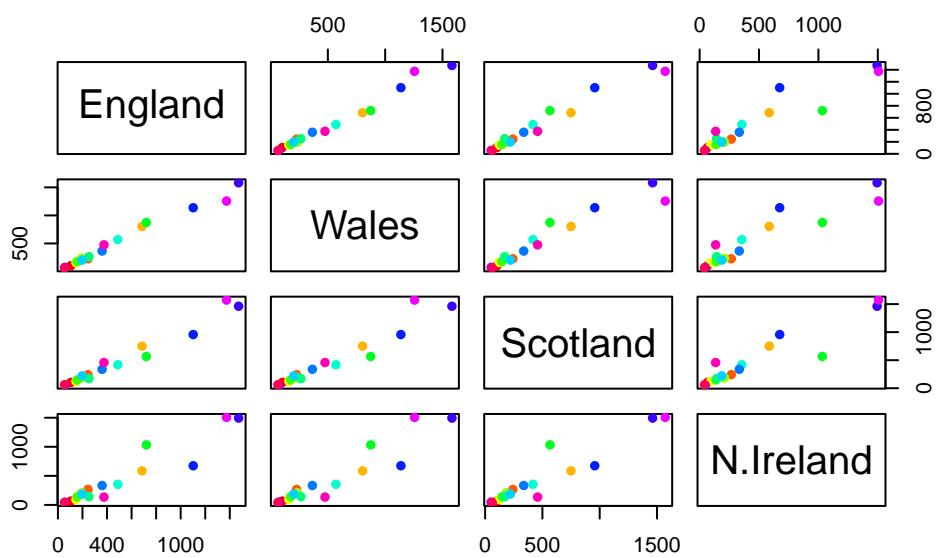


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```

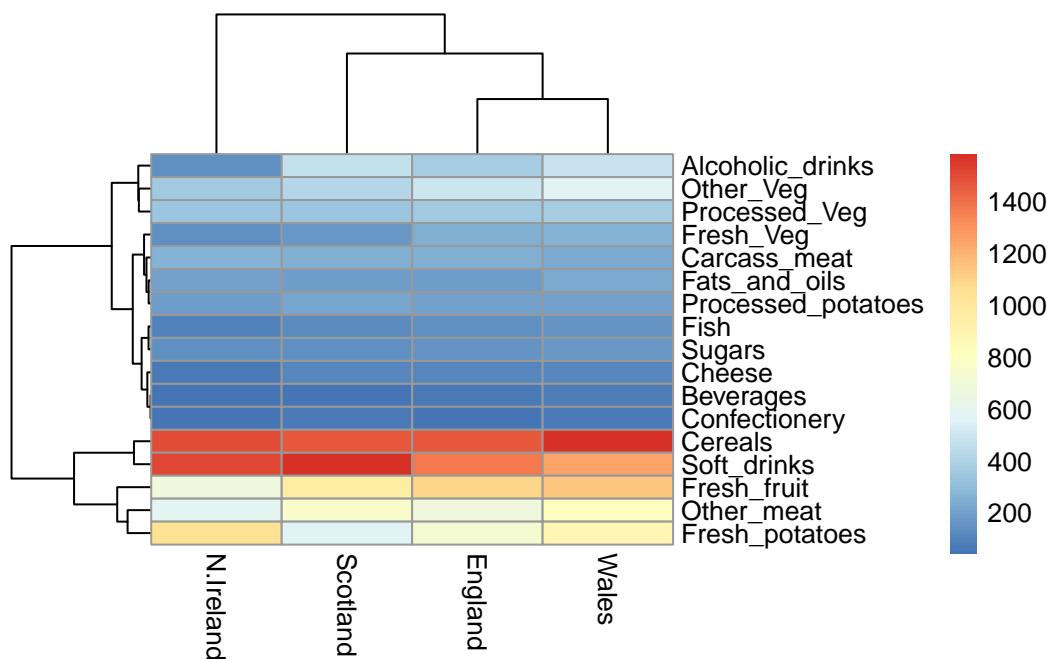


### Pairs plots and heatmaps

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



## PCA to the rescue

The main PCA function in “base R” is called `prcomp()`. This function wants the transpose of our food data as input (i.e. the foods as columns and the countries as rows).

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

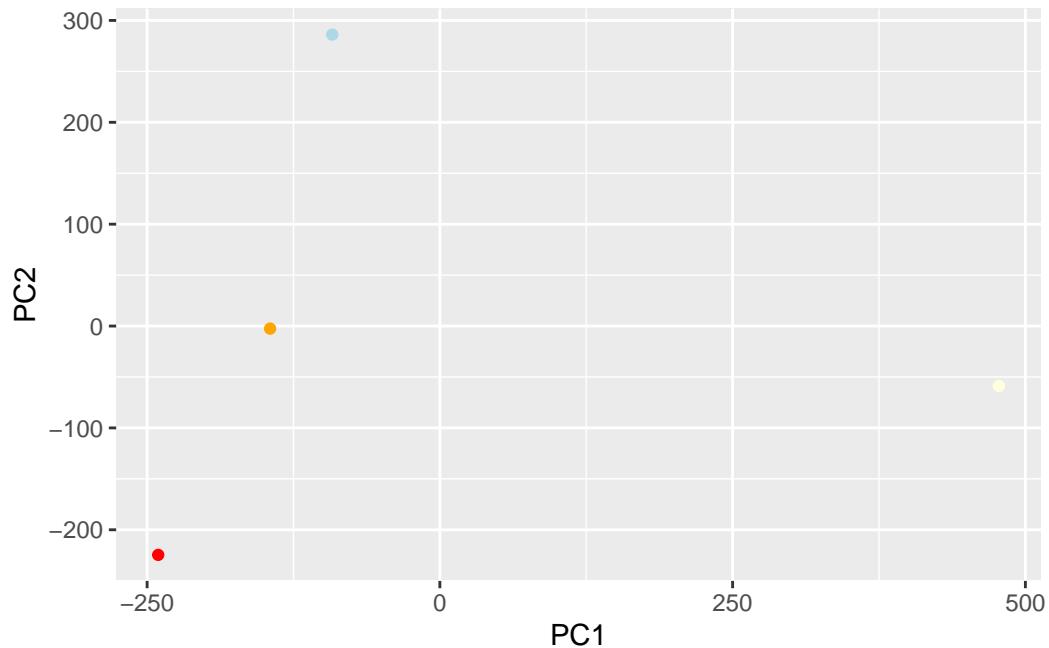
```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation"   "center"    "scale"     "x"

$class
[1] "prcomp"
```

To make one of main PCA result figures we turn to `pca$x` the scores along our new PCs. This is called “PC plot” or “score plot” or “Ordination plot” ...

```
library(ggplot2)
my_cols <- c("orange", "red", "lightblue", "lightyellow")
ggplot(pca$x)+ aes(PC1,PC2, label= rownames(pca$x)) +
  geom_point(col= my_cols)
```



the second major result figure is called a “loadings plot” of “variable contributions plot” or “weight plot”

```
ggplot(pca$rotation) +  
  aes(PC1, rownames(pca$rotation)) +  
  geom_col()
```

